

## 9. Random forests, perceptrons and support vector machines

**Shabana K M**

PhD Research Scholar

Computer Science and Engineering

IIT Palakkad

21 November 2021



INDIAN INSTITUTE  
OF TECHNOLOGY  
**PALAKKAD**



# Recap

- Decision trees
  - classification and regression trees
  - learning a decision tree
  - selecting the splitting attribute
  - information gain, Gini impurity
  - pros and cons
  - avoiding overfitting of trees
  - bias and variance

# Bagging

## Ensemble method


- combines predictions from multiple machine learning algorithms together to make more accurate predictions than any individual model

## Bootstrap Aggregation or Bagging

- simple and very powerful ensemble method
- used to reduce the variance for algorithms with high variance (eg:- decision trees)
- each model in the ensemble built using a bootstrap sample of training data

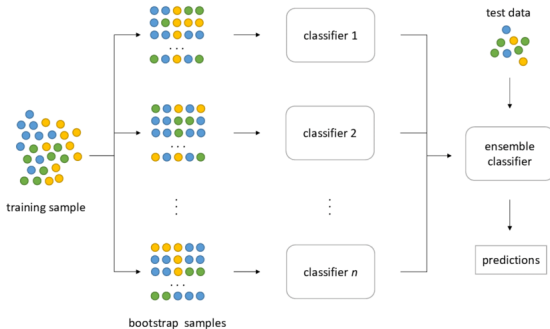
# Bootstrap sampling

- Given a dataset  $D$  containing  $N$  training examples, generate another dataset  $D'$  by drawing  $N$  samples at random with replacement from  $D$



# Bagging

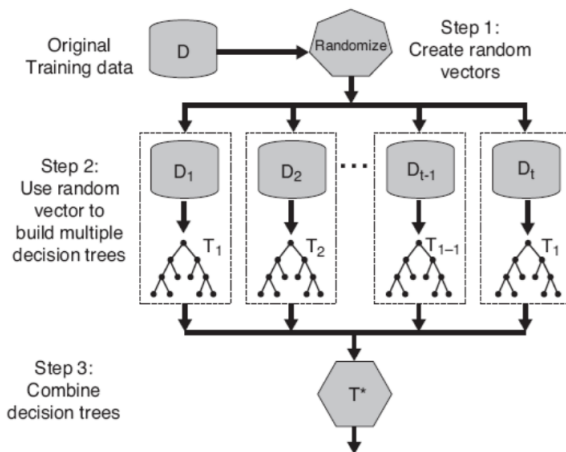
- 1 Create  $k$  bootstrap samples  $D_1, \dots, D_k$
- 2 Train distinct model on each  $D_i$
- 3 Classify new instance by majority vote/average



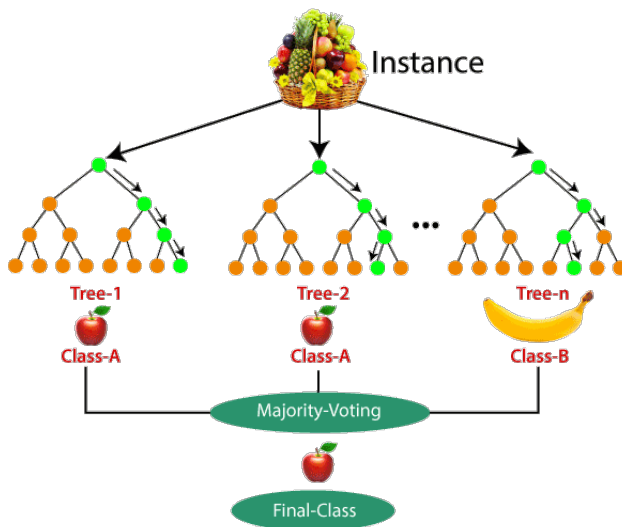
# Random forests

- Grow a forest of many trees
- Each tree grown on an independent **bootstrap sample** from the training data
- At each node:
  - 1 **Random vector method:** select  $m$  features at random out of all possible features (independently for each node)
  - 2 find the best split based on the  $m$  selected features
- Grow the trees to maximum depth (classification)
- Vote/average the trees to get predictions for new data

# Random forests: The idea



# Prediction using random forest






# Perceptrons

- supervised learning algorithm for binary classification
- goal: find a separating hyperplane
  - divide the input space into two classes based on their labels
  - for separable data, guaranteed to find one
- linear classifier
  - makes predictions by combining a set of weights with the feature vector
- online learning algorithm
  - processes one example at a time

## Perceptron as an artificial neuron




- receives input from  $D$  other neurons, one for each input feature
- each incoming connection has a weight
  - feature values are strength of inputs
- neuron sums up all the weighted inputs
- based on this sum decides whether to “fire” or not
  - *firing*: positive example
  - *not-firing*: negative example

# Perceptron: Activation

- The amount of **activation** of a neuron is given by:  $a = \sum_{i=1}^D w_i x_i$
- Let  $b$  be the threshold for activation, i.e., if  $a \geq b$ , the neuron fires
- $\sum_{i=1}^D w_i x_i \geq b$  is equivalent to  $\sum_{i=1}^D w_i x_i - b \geq 0$   
where  $b$  is referred to as **bias**
- Thus if  $w^T x \geq 0$ , then  $x$  is predicted as a positive example  
where  $w = [w_0 \ w_1 \ \dots \ w_D]$ ,  $x = [-1 \ x_1 \ \dots \ x_D]$   
Here  $w_0$  represents the *bias*

# Perceptron: decision boundary



# Perceptron learning algorithm

**Input:** A sequence of training examples  $(x_1, y_1), (x_2, y_2), \dots$  where all  $x_i \in \mathcal{R}^D$ ,  $y_i \in \{-1, 1\}$

- Initialize  $w_0 = 0 \in \mathcal{D}$
- for  $\text{iter}=1, \dots, \text{maxIter}$ 
  - for each training example  $(x_i, y_i)$ :
    - predict  $y' = \text{sign}(w_t^T x_i)$
    - if  $y_i \neq y'$ :  $w_{t+1} = w_t + y_i x_i$

# Perceptron learning algorithm

- **online algorithm:** processes one example at a time
- **error-driven algorithm:** updates parameters only when it makes an error
- weight updates
  - on positive examples:  $w_{t+1} = w_t + x_i$
  - on negative examples:  $w_{t+1} = w_t - x_i$

# Intuition behind the weight update

Suppose the perceptron has made a mistake on a positive example  
i.e.,  $y = 1$  and  $w_t^T x < 0$




**New weight vector**  $w_{t+1} = w_t + x$

**New dot product** will be:  $w_{t+1}^T x = (w_t + x)^T x = w_t^T x + x^T x > w_t^T x$

For a positive example, the perceptron update will increase the score assigned to the same input




Similar reasoning for negative examples

# Geometry of the perceptron update






# Geometry of the perceptron update




# Limitations of perceptron

## 1. Works only for linearly separable data




# Margin

The **margin** of a hyperplane for a dataset is the distance between the hyperplane and the data point nearest to it




# Limitations of perceptron

## 2. Finds any hyperplane separating the two classes




- decision boundary close to training data (unstable)
- prefer a larger margin for generalization

# Which decision boundary to choose?




**maximum margin solution:** most stable under perturbations of the inputs

# Support vectors




- data points that lie closest to the decision surface/hyperplane
- data points most difficult to classify
- have direct bearing on the optimum location of the decision surface

# Support vector machines (SVM)



- maximize the margin around the separating hyperplane
- decision function fully specified by the support vectors

# Support vector machines



- Choose normalization for  $w$  such that  $w^T x_+ + b = +1$  and  $w^T x_- + b = -1$ 
  - $x_+$ : support vector for positive class
  - $x_-$ : support vector for negative class
- **Margin** is given by  $\frac{w}{\|w\|} \cdot (x_+ - x_-) = \frac{w^T x_+ - w^T x_-}{\|w\|} = \frac{2}{\|w\|}$



# SVM - Optimization

- Learning the SVM can be formulated as an optimization problem


$$\max_w \frac{2}{||w||} \text{ subject to } w^T x_i + b \begin{cases} \geq +1 & \text{if } y_i = +1 \\ \leq -1 & \text{if } y_i = -1 \end{cases} \text{ for } i = 1, \dots, N$$

- Or equivalently,


$$\min_w ||w||^2 \text{ subject to } y_i(w^T x_i + b) \geq 1 \text{ for } i = 1, \dots, N$$

- This is a quadratic optimization problem subject to linear constraints and there is a unique minimum

# Linear separability: Which hyperplane?



- the points can be linearly separated but there is a very narrow margin



- but possibly the large margin solution is better, even though one constraint is violated

- In general there is a trade off between the margin and the number of mistakes on the training data

# References

- 1 <https://people.csail.mit.edu/dsontag/courses/ml16/slides/lecture11.pdf>
- 2 <https://www.cs.utah.edu/~zhe/pdf/lec-10-perceptron-upload.pdf>
- 3 [http://ciml.info/dl/v0\\_99/ciml-v0\\_99-ch04.pdf](http://ciml.info/dl/v0_99/ciml-v0_99-ch04.pdf)
- 4 <http://web.mit.edu/6.034/wwwbob/svm-notes-long-08.pdf>
- 5 <https://www.robots.ox.ac.uk/~az/lectures/ml/lect2.pdf>

Thanks Google for the pictures!